

# C for Science

Week 2

Pointers and Dynamic Memory, C-Standard Library Functions and Makefiles

Dr Maria Valera(m.valera-espina@doc.ic.ac.uk)

Imperial College London

February 2016

## Table of Contents week 1

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Course Content

- Program structure.
- C Types: conversions and casts.
- Control structures: sequence,selection and repetition.
- Conditions: operands, operators and its precendence.
- Arithmetic and Logical expressions.
- C Standard Library Functions.
- string Types.
- Pointers and Arrays Types.
- Memory management.
- Input/Output.
- Structures: Type Data Structure and Dynamic Data Stucture.
- C Standard Library and Scientific C-Libraries.
- Optimisation & Debugging.

## Table of Contents

- 1 Recap - Where are we?
- 2 **Mathematical Functions**
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## More Mathematical Functions in `<math.h>`

- Maths functions come with the ANSI Standard C Library, which contains many maths functions. To use them we need a:

- `#include <math.h>`

- Here some example functions:

```
sin(x) asin(x) sinh(x) exp(x)
cos(x) acos(x) cosh(x) log(x)
tan(x) atan(x) tanh(x) log10(x)
sqrt(x) atan2(x,y) pow(x,y) fabs(x)

(all the trigonometric functions use radians!)
```

## The `pow(x, y)` function declared in `<math.h>`

### Exponentiation

As you noticed from your exercise, there is no exponentiation operator (e.g. `^`) in C. Instead, we have the following:

```
xy = pow(x, y)
```

`pow(x, y)` assumes `x` and `y` are of type `double`.

### Notice

The `pow` function is often implemented as:

```
exp(y * ln(x))
```

For whole integer powers (i.e.  $x^2$ ), the multiplication explicitly (`x * x`) should be used.

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 **Pointer Types**
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## This lecture

### Pointers:

- We touched on pointers previously when we looked at `scanf()`.
- Using pointers, we can construct references to *any* C variable.
- Pointers in C are much more powerful, and dangerous. If we use a pointer and the value it points to no longer exists, **bad** things will happen.

### Dynamic memory allocation:

- We'll be looking at dynamic (at run-time) memory allocation. <sup>1</sup>
- Failure of code to do this is called a *memory leak*.

**Incorrect use of pointers and memory handling are major sources of C programming errors.**

<sup>1</sup>C's equivalent of the **new** operator in Java.

## Pointers

- Memory can be seen as an ordered sequence of consecutively numbered storage locations.
  - Variables are stored in memory in one or more adjacent storage locations depending on its type:

```
char str[] = "is"; int i = 1; char c = 'R';
```

i	s	\0	1	R
0xbffff72	0xbffff73	0xbffff74	0xbffff75	0xbffff76
str			i	c

- The **address** (&) of a variable is the address of its first storage location.

i	s	\0	1	R
0xbffff72	0xbffff73	0xbffff74	0xbffff75	0xbffff76
&str==0xbffff72			&i==0xbffff75	
			&c==0xbffff79	

- A pointer in C represents the starting address of a value in memory.

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax**
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Declaring pointers

- The syntax for declaring pointers in C can be slightly inconsistent.
- For primitive types, we simply append a "\*" to the type to construct the pointer type:

### Declare a pointer

```
double* doublePtr; /* A pointer to a double */
int *intptr; /* A pointer to an int */
char* charptr, strptr; /* A pointer to a char 'charptr'
                        and a pointer to a char 'strptr' */
int **intPtrPtr; /* A pointer to a pointer to an int */
```

- After declaring a pointer we assign the first storage location of a variable to a pointer using &

```
strptr = (char *) &str; intptr = &i; charptr = &c;
```

0xbffff72	0xbffff75	0xbffff79
0xbffff80	0xbffff81	0xbffff82
0xbffff83	0xbffff84	0xbffff85
strptr	intptr	charptr

## The const keyword

The `const` keyword allows us to specify that certain values cannot be modified.

```
const int x = 4;
x = 5; /* This is a compile-time error */
```

`const` can also be applied to pointers, though it can become confusing quickly for multiple levels of indirection:

```
int val = 5;
const int *ptr1 = &val; /* ptr1 can be modified, val cannot */
int *const ptr2 = &val; /* val can be modified, ptr2 cannot */
const int *const ptr3 = &val; /* neither val nor ptr3 can be modified */
```

Try reading the pointer declarations right to left.

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operators**
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Pointer Operators


The following two operators are the primary mechanism for performing pointer-related operations in C:

- The address operator (`&`) is a prefix operator that takes a value and returns its address.
- The indirection operator (`*`) is a prefix operator that takes a pointer and returns the value it points to. This is called “de-referencing” the pointer.

## Pointer Operators


### Declare and Assign

```
int x = 42;
int *intPtr = &x;
```



### De-referencing

```
*intPtr = 77;
```



## Printing pointers

`printf` can also print pointer values using the “%p” specifier. Printing out the addresses of a parameter in a recursive function lets us see what way the stack grows:

```
#include <stdio.h>

void printAddresses(int depth) {
    printf("Address of depth: %p\n", &depth);

    if(depth > 0) {
        printAddresses(depth-1);
    }
}

int main(void) {
    printAddresses(5);
    return 0;
}
```

## A word on null pointers

- You can nullify any pointer by assigning the integer `0` to it.
- The header file `stdlib.h` defines the macro `NULL` which is equivalent.
- Checking if a pointer is `NULL` is a good thing to do, but only useful if it was actually set to `NULL` in the first place.

**intPtr might not be NULL!**

```
#include <stdlib.h>
#include <assert.h>

void someFunction(int *output) {
    assert(output != NULL);
    *output = 42;
}

int main(void) {
    int *intPtr;
    someFunction(intPtr);
    return 0;
}
```

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic**
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Pointer Arithmetic

In C, we have the ability to explicitly manipulate pointer values. We can add and subtract values from them, and use the increment (`++`) and decrement (`--`) operators on them.

```
#include <stdio.h>

int main(void) {
    char greeting[] = "Hello world!\n";
    char *currentLetter = greeting;

    while(*currentLetter != '\0') {
        putchar(*currentLetter);
        ++currentLetter;
    }

    return 0;
}
```

## Pointer Arithmetic

Incrementing and decrementing pointers does so in multiples of the size of the type being pointed to:

```
#include <stdio.h>

int main(void) {
    char strArr[] = "abcd";
    int intArr[] = {1, 2, 3, 4};

    char *strPtr = strArr;
    int *intPtr = intArr;
    for(int i=0; i<4; ++i) {
        printf("strPtr = %c, strPtr = %p\n", *strPtr, strPtr);
        printf("intPtr = %i, intPtr = %p\n", *intPtr, intPtr);

        ++strPtr;
        ++intPtr;
    }
}
```

## Pointer Arithmetic


```
char strArr[] = "abcd";
int intArr[] = {1, 2, 3, 4};
```

```
char *strPtr = strArr;
int *intPtr = intArr;
```

```
strPtr++;
intPtr++;
```

```
intPtr--;
```

```
/* Circumventing the type system - don't do this! */
intPtr = (int*) (((char*) intPtr) + 1);
```



## Allowed Pointer Operations

- Declaration: `double *pA, *pB;`
- Assignment: `pA = &var;`
- Increment: `pA = pA + 1; /*Incrementing memory position */`
- Decrement: `pA = pA - 1; /*Decrementing memory position */`
- Difference: `gap = pA - pB; /*Subtraction is a block of memory */`
- Comparison: `if(pA == pB) /* comparison of memory positions */`
- De-referencing: `*pA = val; *pA = *pA + 1; /*Incrementing var variable to value + 1*/`

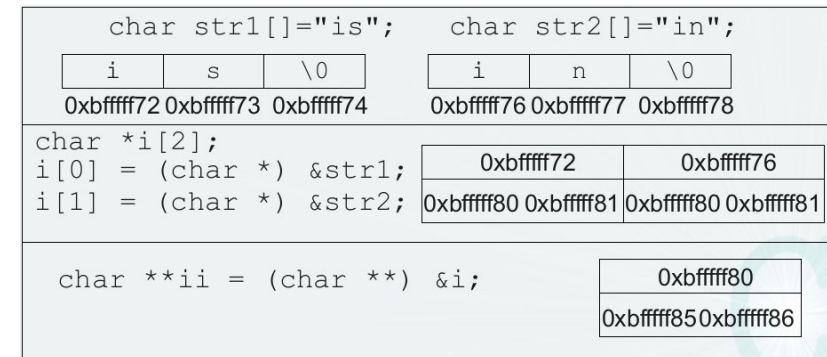
## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 **Pointer to Pointer**
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Pointer to Pointer

```
#include <stdio.h>
int main(){
    char str1[]="is"; char str2[]="in";
    char *i[2];
    i[0] = (char *) &str1;
    i[1] = (char *) &str2;
    char **ii = (char **) &i;
    printf("will print 'is': %s\n",i[0]);
    printf("will print 'in': %s\n",i[1]);
    printf("will print 'is': %s\n",*ii);
    ii++;
    printf("will print 'in': %s\n",*ii);
    return 0;
}
```

## Pointer to Pointer



## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers**
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Pointers and Arrays

- The concepts of arrays and pointers are closely related in C.
- An array can be casted *implicitly* to a pointer to the element type.

```
double x[10];
double *xPtr = x;
float *floatPtrArray[10]; /* An array of pointers to floats */
int y[15][30];
int (*yPtr)[30] = y; /* A pointer to an array of floats */
```

- What happens if we want a pointer to the first element of a multi-dimensional array in C?

```
float array2D[10][20];
float *array2DPtr = &array2D[0][0];
```

- We simply take the address of the first element.

## Pointers using Array Notation

Rather than using pointer arithmetic and the dereference operator, it is possible to use the `[]` operator in exactly the same way as with arrays:

The following three accesses are identical:

```
void func(char *str) {
    /* Access the eighth letter (we count from 0) */
    char letter1 = *(str+7);
    char letter2 = str[7];
    char letter3 = (str+4)[3];
}
```

## Fixed Size Two-Dimensional Arrays

As we saw in week2, we can declare arrays of dimension higher than one, as follows:

```
double a[2][3] = {{1.0, 2.0, 3.0},{2.0, 3.0, 4.0}};
```

Where the elements of `a` are denoted as:

<code>a[0][0]</code>	<code>a[0][1]</code>	<code>a[0][2]</code>
<code>a[1][0]</code>	<code>a[1][1]</code>	<code>a[1][2]</code>

In memory it is arranged as follows:

<code>a[0][0]</code>	<code>a[0][1]</code>	<code>a[0][2]</code>	<code>a[1][0]</code>	<code>a[1][1]</code>	<code>a[1][2]</code>
----------------------	----------------------	----------------------	----------------------	----------------------	----------------------

They are allocated from the stack thus large arrays may cause problems.

To access the top left element:

```
myVal = a[0][0]; /* equal to 1.0 */
```

## Example 2D Array

```
#include <stdio.h>
#define COLS 3

void printArray(int matrix[][COLS], int rows)
{
    int i, j;
    for (i = 0; i < rows; i++) {
        for (j = 0; j < COLS; j++){printf("%d ", matrix[i][j]);}
        printf("\n"); }
}

int main()
{
    int matrix[2][COLS] = {{1, 2, 3},{4, 5, 6}};
    printArray(matrix, 2);
    return 0;
}
```

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary



## Passing by reference

- The pointer passed by reference is copied but not the data pointed to.

```
#include <stdio.h>

void swap(int *a, int *b) {
    int temp = *b;
    *b = *a;
    *a = temp;
}

int main(void) {
    int a = 42;
    int b = 77;

    printf("a: %i, b: %i\n", a, b);
    swap(&a, &b);
    printf("a: %i, b: %i\n", a, b);

    return 0;
}
```

```
#include <stdio.h>

void swap(int *a, int *b);

int main (void){
    int a =3;
    int b = 4;
    printf("a: %i , b: %i\n",a,b);
    swap(&a,&b);
    printf("a: %i , b: %i\n",a,b);
    return 0;
}

void swap (int *a,int *b){
    int *temp = b;
    *b = *a;
    *a = *temp;
}
```

## The truth about passing arrays to functions

- C doesn't perform type checking on the leading dimension of an array when passing it to a function.
- C passes arrays by reference.
- This may have seemed a little strange given that the only way to pass by reference in C is to use a pointer.
- In fact, C always converts the leading dimension of an array to a pointer when passing it to a function.
- This is the true reason why arrays are passed by reference, and why the leading dimension is never checked.

## The truth about passing arrays to functions

We can see the artifacts of this behaviour if we look close enough. What do you think this will print?

```
#include <stdio.h>

void printSizeArray(char array[100]) {
    printf("sizeof(array) = %li\n", sizeof(array));
}

void printSizePtr(char *ptr) {
    printf("sizeof(ptr) = %li\n", sizeof(ptr));
}

int main(void) {
    char buffer[100];
    printf("sizeof(buffer) = %li\n", sizeof(buffer));
    printSizeArray(buffer);
    printSizePtr(buffer);
    return 0;
}
```

## Command Line Arguments

The main function of a C program can also have a type signature where it receives arguments passed to it from the command line. It takes two parameters:

- argc - The number of parameters passed.
- argv - An array of C strings.

We can print them as follows:

[Prints each argument](#)

```
#include <stdio.h>

int main(int argc, char **argv) {
    for(int i=0; i<argc; ++i) {
        printf("argv[%i] = %s\n", i, argv[i]);
    }

    return 0;
}
```

## atoi()

You will need to convert a command line argument to a number for your tutorials. The easiest way to do this is with the `atoi()` function from `stdlib.h`.

```
int atoi(const char *nptr);/* atoi() does not detect errors*/
```

Counts from one to the supplied argument

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv){
    int count =0;
    if(argc >1) {
        count = atoi(argv[1]);
    }
    printf("the number of count is %d\n", count);
    for(int i=0; i< count; i++){
        printf("%d\n",i);
    }
}
```

## Pointers as function types

- We also can declare pointers as function types:

```
char *doSomething()
```

- `doSomething()` will return a `char` pointer. However, this is only really useful if dynamic memory is allocated.

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions**
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## String Utilities

The header `string.h` contains a number of useful utility functions:

- String length:

```
size_t strlen(const char *s);
```

- String concatenation:

Requires `dest` to have size `strlen(dest)+n+1!`

```
char *strncat(char *dest, const char *src, size_t n);
```

- String comparison:

```
int strncmp(const char *s1, const char *s2, size_t n);
```

- String copying:

Both unsafe!

```
char *strcpy(char *dest, const char *src);
char *strncpy(char *dest, const char *src, size_t n);
```

## On Unsafe String Functions

The C API contains many unsafe string functions. Whenever you use them, check the documentation for the following:

- Whenever a buffer is being written to, the requirements on the size of the destination buffer.
- If the function writes a string, under what circumstances it terminates the string with a `'\0'`.

If the source string is long enough, `strncpy` will not terminate the destination string with `'\0'`, possibly causing later code to run off the end of the string. This is why it is considered unsafe.

Have a look at `strncpy()` and `strlcat()` for an example of safer functions. They both come from BSD, and are therefore unfortunately non-portable.

## String Utilities

String comparison:

```
int strcmp(const char *s1, const char *s2);
```

Lexicographically compares the strings `s1` and `s2`. The return value is

- `<0` - if `str1` is less than `str2`.
- `=0` - if `str1` is equal to `str2`.
- `>0` - if `str1` is greater than `str2`.

There is also a function `strncmp` which only compares up to `n` characters:

```
int strncmp(const char *s1, const char *s2, size_t n);
```

## strcmp() example

```
#include <stdio.h>
#include <string.h>

int main(void) {
    char a[] = "astring";
    char b[] = "astring";
    char c[] = "astr ing";

    if (!strcmp(a, b)) {
        printf("Strings a and b are the same\n");
    }

    if (!strcmp(a, c)) {
        printf("Strings a and c are the same\n");
    }

    return 0;
}
```

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers**
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Function Pointers

C also supports pointers to functions. Here's how we can take a pointer to a sum function:

```
static int sum(int a, int b) {
    return a + b;
}

int main(void) {
    int (*sum_ptr)(int, int);
    sum_ptr = &sum;
    return 0;
}
```

We've written the declaration of `sum_ptr` the same way we'd have written a function declaration except we replaced the function name with `(*sum_ptr)`.

## Function Pointers

It's possible to invoke a function pointer in exactly the same way as normal function.

```
#include <stdio.h>

static int sum(int a, int b) {
    return a + b;
}

int main(void) {
    int (*sum_ptr)(int, int) = &sum;

    printf("The sum of 39 and 73 is %i.\n", sum_ptr(39, 73));
    return 0;
}
```

## Function Pointers

We can pass them to other functions as well.

```
#include <stdio.h>

static int sum(int a, int b) { return a + b; }
static int mul(int a, int b) { return a * b; }

static void print_result(int (*func)(int, int), int a, int b) {
    printf("func(%i, %i) = %i\n", a, b, func(a, b));
}

int main(void) {
    int a = 42;
    int b = 37;

    print_result(&sum, a, b);
    print_result(&mul, a, b);
    return 0;
}
```

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Dynamic Memory

- Up until this point, we've only used stack allocated values – values which are destroyed as soon as they go out of scope.
- There is another area of memory called the *heap* which can hold dynamic memory.
- In Java, destroying unreferenced dynamically allocated values was done through a process of *garbage collection*.
- In C, you will need to design your own strategy for freeing dynamic memory, depending on the context.

## malloc() example

### A C program

```
#include <stdlib.h>
#include <stdio.h>

int main(void) {
    int size = 100;
    int *arr = malloc(size * sizeof(int));

    for(int i=0; i<size; ++i){
        arr[i] = i;
    }
    free(arr);
    return EXIT_SUCCESS;
}
```

### A C equivalent

- &arr - The address of the arr in the stack
- \*arr - The contents of what is in the heap
- arr - The address allocated in the heap

## malloc() example

- malloc() is completely unaware of how the returned memory will be used.
- malloc takes its size parameter in bytes, *not elements*. We need to use `sizeof()` to work out how much memory to use.
- malloc returns a type of `void*` (a pointer to an unknown type). Any value pointer can be implicitly converted to `void*`, but `void*` must be explicitly cast to another pointer type.

### We could have written:

```
void *memory = malloc(size * sizeof(int));
int *arr = (int*) memory;
```

## Using dynamic memory safely

- malloc(), realloc() and calloc() all return NULL if the allocation fails. This is *not* a fatal error, so you must check for it.
- The pointer passed to free() *must* come from malloc (or its relatives) or be NULL.
- Except for calloc(), the memory allocation routines return uninitialised memory.
- If you can't find the free() corresponding to a memory allocation in your code, you may have a bug. It won't cause your code to crash (unless you exhaust RAM) but can cause your program's memory use to become bloated.
- **Remember to free!**

## The Dynamic Memory API

The main allocation-related functions in `stdlib.h` are:

- `malloc()` - allocates a region of memory of size bytes and returns a pointer to the allocated memory.

```
void *malloc(size_t size);
```

- `calloc()` - allocates a region of memory that can hold `nmemb` elements of size bytes. The region is initialised to 0.

```
void *calloc(size_t nmemb, size_t size);
```

- `realloc()` - reallocates a region of memory to the supplied size, preserving the contents.

```
void *realloc(void *ptr, size_t size);
```

- `free()` - frees a memory region previously allocated using the above.

```
void free(void *ptr);
```

## When `malloc()` fails

- In the code you write, you probably don't have much choice except to exit if a `malloc()` fails. This might not be acceptable in other cases (e.g. a kernel).
- The `perror()` function defined in `stdio.h` prints the last error encountered by a system or library routine to standard error, prefixed by the supplied string (allowed to be NULL).
- The `exit()` function defined in `stdlib.h` immediately (but relatively cleanly) terminates the process with the supplied status code.
- `EXIT_SUCCESS` and `EXIT_FAILURE` are error codes defined in `stdlib.h` and are slightly more portable to non-POSIX systems than using 0 and non-zero values for exit codes.

## A Quick Recap

After executing this code:

```
/* Allocates space for two integers */
int *intPtr = malloc(sizeof(int) * 2);
```

- `&intPtr` is the address of the pointer on the stack, and has the type `int**`.
- `intPtr` is a stack-allocated value which contains the starting address of the region allocated on the heap and has the type `int *`.
- `*intPtr` or `intPtr[0]` is the value (uninitialised) of the first integer in the heap-allocated region.
- `*(intPtr+1)` or `intPtr[1]` is the value (uninitialised) of the second integer in the heap-allocated region.

## Clearing Memory

Since both stack and heap allocated memory may contain uninitialised values, it's useful to be able to zero large regions quickly. We can do this with the `memset()` method from `string.h`.

```
void *memset(void *s, int c, size_t n);
```

The `n`-byte region pointed to by `s` is set to the value `c`. Although `c` is an `int`, it is converted to an `unsigned char` first. `s` is returned.

`memset()` example

```
char quote[] = "To be or not to be";
memset(quote, '.', 9);
printf("%s\n", quote);
```

Output

```
.....not to be
```

## Copying Memory

Copying regions of memory may be done using the `memcpy()` method from `string.h`.

```
void *memcpy(void *dest, const void *src, size_t n);
```

Copies `n` bytes from `src` to `dst`, returning `dest`. The source and destination regions must not overlap. If they do, use `memmove`.

### `memcpy()` example

```
char str[] = "Morning World!";
char time[] = "Evening";
memcpy(str, time, 7);
printf("%s\n", str);
```

### Output

Evening World!

## Valgrind

- Valgrind is a GPL-licensed framework for debugging and profiling tools that can run under Linux and Mac OS.
- It functions by disassembling the application at run-time, adding instrumentation instructions and then converting back to machine code.
- You can expect Valgrind to result in a slowdown of around 5-100x.

## Valgrind

Valgrind comes with a number of tools that use the core framework. A few of them are:

**Memcheck** Detected invalid memory accesses, use of uninitialised memory, memory leaks, invalid uses of free and other errors.

**Callgrind** Provides detailed call-graph information, and with the “-simulate-cache” option, estimated values of cache hits/misses and cycle counts.

**Massif** Produces information on the heap usage of a program.

**Helgrind** Locates data races in multi-threaded programs. Specifically, it locates values that are accessed by multiple threads that do not appear to have an associated lock.

## Memcheck

- Invoking memcheck:

```
$ valgrind --tool=memcheck ./executable
```

- If we compile with the `-g` option to the C compiler, Valgrind will give us line numbers.
- Supplying `--leak-check=full` to Valgrind will give details of individual memory leaks.

## An example

The following example overruns the bounds of the heap-allocated region and fails to free it afterwards.

```
#include <stdlib.h>

int main(void)
{
    double *squares = malloc(100 * sizeof(double));

    for(int i = 0; i <= 100; ++i)
        squares[i] = i * i;

    return EXIT_SUCCESS;
}
```

## Example Memcheck output

```
==974== Memcheck, a memory error detector
==974== Copyright (C) 2002-2011, and GNU GPL'd, by Julian Seward et al.
==974== Using Valgrind-3.7.0 and LibVEX; rerun with -h for copyright info
==974== Command: ./broken_memory
==974==
==974== Invalid write of size 8
==974==    at 0x40055A: main (broken_memory.c:8)
==974== Address 0x51f1360 is 0 bytes after a block of size 800 alloc'd
==974==    at 0x4C2B3F8: malloc (in /usr/lib/valgrind/vgpreload_memcheck-amd64-linux.so)
==974==    by 0x40052D: main (broken_memory.c:5)
==974==
==974== HEAP SUMMARY:
==974==    in use at exit: 800 bytes in 1 blocks
==974==    total heap usage: 1 allocs, 0 frees, 800 bytes allocated
==974==
==974== 800 bytes in 1 blocks are definitely lost in loss record 1 of 1
==974==    at 0x4C2B3F8: malloc (in /usr/lib/valgrind/vgpreload_memcheck-amd64-linux.so)
==974==    by 0x40052D: main (broken_memory.c:5)
==974==
==974== LEAK SUMMARY:
==974==    definitely lost: 800 bytes in 1 blocks
==974==    indirectly lost: 0 bytes in 0 blocks
==974==    possibly lost: 0 bytes in 0 blocks
==974==    still reachable: 0 bytes in 0 blocks
==974==    suppressed: 0 bytes in 0 blocks
==974==
==974== For counts of detected and suppressed errors, rerun with: -v
==974== ERROR SUMMARY: 2 errors from 2 contexts (suppressed: 2 from 2)
```

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Constructing Matrices with Pointers

### Allocate Dynamic Memory

```
double** makeMatrix(unsigned int rows, unsigned int cols)
{
    unsigned int i;
    double** matrix;

    matrix = (double**) malloc(rows * sizeof(double *));
    if (!matrix) { return NULL; } /* failed */

    for (i = 0; i < rows; i++)
    {
        matrix[i] = (double *) malloc(cols*sizeof(double));
        if (!matrix[i])
            return NULL; /* lazy, we should really free
                           all the memory allocated above */
    }

    return matrix;
}
```



## Accessing Matrix Elements

### Usage pattern for makeMatrix

```
double** matrix = makeMatrix(rows, cols);
for (i=0; i < rows; i++){
    for (j=0; j < cols; j++){
        matrix[i][j] = 0.0;
    }
}
//free the matrix
```

- Accessing the dynamically allocated array looks identical to the fixed size ones, but “under the hood” things are a little different:  
`matrix[row][col] = (*(matrix + row) + col)`
- The `makeMatrix` code on the previous slide contained a lot of `malloc` statements, is there a better way to allocate a matrix?

## Another way of Allocating Matrices

### Allocate Dynamic Memory

```
double** makeMatrix(unsigned int rows, unsigned int cols)
{
    unsigned int i;
    double** matrix;

    matrix = (double**) malloc(rows * sizeof(double *));
    if (!matrix) { return NULL; } /* failed */

    for (i = 0; i < rows; i++)
    {
        matrix[i] = (double *) malloc(cols*sizeof(double));
        if (!matrix[i])
            return NULL; /* lazy, we should really free
                           all the memory allocated above */
    }

    return matrix;
}
```

## Why is allocMatrix better?

- `allocMatrix` only uses 2 mallocs whilst, `makeMatrix` uses `cols + 1`.
- Meaning there are fewer points of failure (we only check two pointers for NULL).
- It is much easier to free a matrix allocated with the `allocMatrix` function, all we need to do is:

### Free Dynamic Memory

```
void freeMatrix(double** matrix)
{
    free(matrix[0]);
    free(matrix);
}
```

## Matrices utility functions

Let's define some utility functions to:

- Allocate memory for the matrix (`allocMatrix`) - done.
- Free a matrix (`freeMatrix`) - done.
- Print a matrix (`printMatrix`).
- Create a random matrix (`randomMatrix`).
- Add matrices together (`addMatrix`).

## printMatrix, randomMatrix and addMatrix

### printMatrix

```
void printMatrix(double** matrix, unsigned int rows,
                unsigned int cols)
{
    unsigned int i, j;
    for (i = 0; i < rows; i++){
        for (j = 0; j < cols; j++){
            printf("%8.5lf ", matrix[i][j]);
        }
        printf("\n");
    }
}
```

### randomMatrix

```
void randomMatrix(double** matrix, unsigned int rows,
                 unsigned int cols)
{
    unsigned int i, j;
    for (i = 0; i < rows; i++){
        for (j = 0; j < cols; j++){
            matrix[i][j] = (double)rand()/RAND_MAX;
        }
    }
}
```

### addMatrix

```
void addMatrices(double** matrixA, double** matrixB,
                double** matrixR, unsigned int rows,
                unsigned int cols)
{
    unsigned int i, j;
    for (i = 0; i < rows; i++){
        for (j = 0; j < cols; j++){
            matrixR[i][j] = matrixA[i][j] + matrixB[i][j];
        }
    }
}
```

## The main function

### main not completed..

```
int main(void)
{
    unsigned int rows, cols;
    double ** matrixA, ** matrixB, **matrixC;
    printf("Enter rows cols: ");
    scanf("%u %u", &rows, &cols);

    matrixA = allocMatrix(rows, cols);
    matrixB = allocMatrix(rows, cols);
    matrixC = allocMatrix(rows, cols);

    if (!matrixA || !matrixB || !matrixC)
    { /* a little lazy, but it does the job */
        fprintf(stderr, "Unable to allocate matrices!\n");
        return -1;
    }

    randomMatrix(matrixA, rows, cols); randomMatrix(matrixB, rows, cols);
    addMatrices(matrixA, matrixB, matrixC, rows, cols);

    printf("\n\nmatrix A = \n");
    printMatrix(matrixA, rows, cols);
    printf("\n\nmatrix B = \n");
    printMatrix(matrixB, rows, cols);
    printf("\n\nmatrix A + matrix B = \n");
    printMatrix(matrixC, rows, cols);

    freeMatrix(matrixC); freeMatrix(matrixB); freeMatrix(matrixA);
}
```

## Results

### Output

```
Enter rows cols: 4 4

matrix A =
0.84019 0.39438 0.78310 0.79844
0.91165 0.19755 0.33522 0.76823
0.27777 0.55397 0.47740 0.62887
0.36478 0.51340 0.95223 0.91620

matrixB =
0.63571 0.71730 0.14160 0.60697
0.01630 0.24289 0.13723 0.80418
0.15668 0.40094 0.12979 0.10881
0.99892 0.21826 0.51293 0.83911

matrixA + matrixB =
1.47590 1.11168 0.92470 1.40541
0.92795 0.44044 0.47245 1.57241
0.43445 0.95491 0.60719 0.73768
1.36371 0.73166 1.46516 1.75531
```

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Header files

- `stdio.h` is one of a number of *header* files defined by the C standard library.
- On Unix-like systems, you can usually find it at the location `/usr/include/stdio.h`. You can open it in a text editor.
- Header files contain information about functions, types and global variables that a library (or other C source files) want to export.
- Header files use exactly the same syntax as C source files, only the “.h” extension distinguishes them as header files.
- We'll look at defining our own headers later.

## Splitting Code Across Multiple Files

- Until now, we've only considered how to write programs whose source code resides in a single file.
- When a program is split across multiple files, the compiler compiles each source file *independently*.
- *Headers* provide just enough information about available functions and variables to type-check.
- Care must be taken to avoid accidentally creating duplicate symbols (names of functions or variables).

## A simple example

`add.h`

```
#ifndef ADD_H
#define ADD_H

int add(int a, int b);

#endif
```

`add.c`

```
#include "add.h"

int add(int a, int b) {
    return a + b;
}
```

`sum.c`

```
#include "add.h"

int main(void) {
    int sum = add(5, 4);
    return 0;
}
```

## Compiling Multiple Source Files

If we assume `add.h`, `add.c` and `sum.c` are in the current directory, we can compile our program as follows:

```
$ gcc -Wall -pedantic sum.c add.c -o sum
```

If we only modify `add.c`, we don't want to have to recompile everything. We can instruct the compiler to create object files (.o files) with the `-c` option.

```
$ gcc -Wall -pedantic -c sum.c -o sum.o
```

```
$ gcc -Wall -pedantic -c add.c -o add.o
```

These can then be combined into the final executable through a process called *linking*:

```
$ gcc sum.o add.o -o sum
```

## Including Headers

- The `#include "file.h"` directive searches the source file directory first, before looking at other paths.
- The `#include <file.h>` only looks at predefined paths and paths given to the compiler on the command line.
- All headers should be surrounded by *include guards*. They have the form:

```
#ifndef __SOME_UNIQUE_TOKEN__
#define __SOME_UNIQUE_TOKEN__
/* Your code */
#endif
```

- This prevents the header file content from being included more than once.

## Table of Contents

- 1 Recap - Where are we?
- 2 Mathematical Functions
- 3 Pointer Types
- 4 Declare Pointer Syntax
- 5 Pointer Operatos
- 6 Pointer Arithmetic
- 7 Pointer to Pointer
- 8 Arrays using pointers
- 9 Pointers as function arguments or function type
- 10 String C-Standard Library Functions
- 11 Function to Pointers
- 12 Memory management
- 13 Allocate 2D arrays dynamically - Example
- 14 Makefile
- 15 Summary

## Summary

In this set of slides, we looked at:

- C's pointer types.
- Using the address (&), indirection (\*) and array subscript ([]) operators with pointers and values.
- Using arithmetic with pointers.
- Why C arrays are passed by reference.
- The command line arguments.
- String Utilities.
- File Manipulation.
- Makefiles.
- Dynamic memory allocation.
- `perror()` and `exit()`.
- `memset()` and `memcpy()`.